

Raptors: NumPy-Compatible Compute for Async Python Services

Raptors is a Rust-backed numerical computing engine that brings NumPy-style computation into modern async Python services without blocking the event loop.

The Problem

NumPy is synchronous and blocks async servers. Developers rely on fragile thread pools, unpredictable CPU usage, and ad-hoc workarounds.

The Solution

Raptors provides async job APIs backed by a Rust execution engine with explicit CPU scheduling, predictable performance, and NumPy-compatible ergonomics.

Why It Wins

- Async-safe by design
- Rust-powered parallelism
- Predictable latency under load
- Drop-in adoption for services

Target Users

Backend engineers, platform teams, and ML infrastructure teams running CPU-bound math inside async Python services.

Positioning

Raptors is not async math and not a NumPy replacement. It is a service-oriented compute engine for production Python systems.